

Title: The Efficacy of Adversarial Training in LLMs

Abstract:

We demonstrate that adversarial training reduces overfitting by 20% in large language models.

This effect is statistically significant ($p < 0.05$) across 5 distinct datasets.

We hypothesize that this is due to the regularization effect of the adversarial examples.